

BIA-667-A: Introduction to Deep Learning and Business Applications

References:

deeplearningbook.org

Dive into Deep Learning, <https://d2l.ai/index.html>

introtodeeplearning.com

deeplearningindaba.com

Agenda


- Introduction
 - Tell us about yourself, e.g. major, interest, expectation from this class, ...
- Syllabus
- Hands-on Labs
 - Try it yourself!
 - Recitation schedule
 - Homework assignments will be similar to labs
- Poll Everywhere
- Introduction to Deep Learning
- Review of Probability Theory

What is Deep Learning?

- A. If sender contains "xxx.com" -> move to Junk
- B. Detect spam emails based on email content
- C. Use Recurrent Neural Network to detect spam emails

ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior




MACHINE LEARNING

Ability to learn without explicitly being programmed




DEEP LEARNING

Extract patterns from data using neural networks



Historical Waves



Three historical waves:

- Cybernetics (1940s–1960s): development of theories of biological learning and implementation of the first models such as the perceptron allowing the training of a single neuron
- Connectionism or Neural Networks (1980–1995): back-propagation to train a neural network with one or two hidden layers
- Deep learning (2006 -)

Number of Neurons





Figure 1.11

(Goodfellow 2016)

ImageNet Challenge: Classification Task



2012: AlexNet. First CNN to win.

- 8 layers, 61 million parameters

2013: ZFNet

- 8 layers, more filters

2014: VGG

- 19 layers

2014: GoogLeNet

- “Inception” modules
- 22 layers, 5 million parameters

2015: ResNet

- 152 layers


Why Deep Learning and Why Now?

Why Deep Learning?

Hand engineered features are time consuming, brittle and not scalable in practice

Can we learn the **underlying features** directly from data?

Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

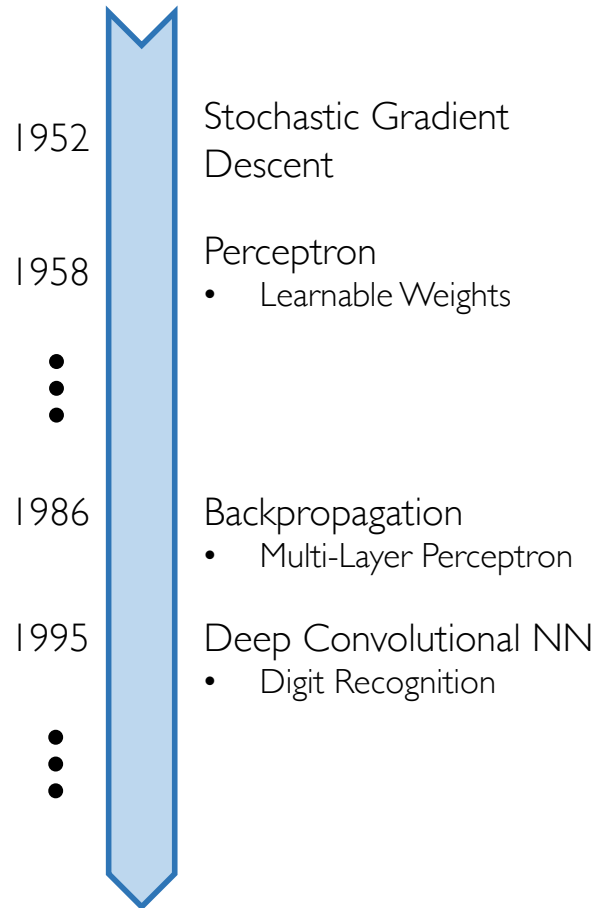
High Level Features



Facial Structure

Why Now?

Neural Networks date back decades, so why the resurgence?



1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



3. Software

- Improved Techniques
- New Models
- Toolboxes



Historical Trends: Growing Datasets




Figure 1.8

Kinds of Deep Learning Problems



Regression: How much a house costs?



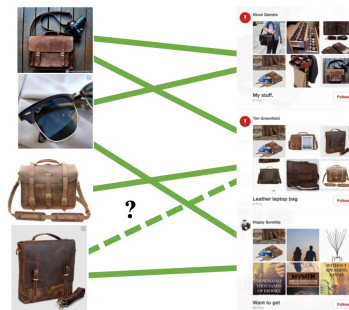
Fig. 1.3.2 Death cap—do not eat!

Classification: Is mushroom poisonous?

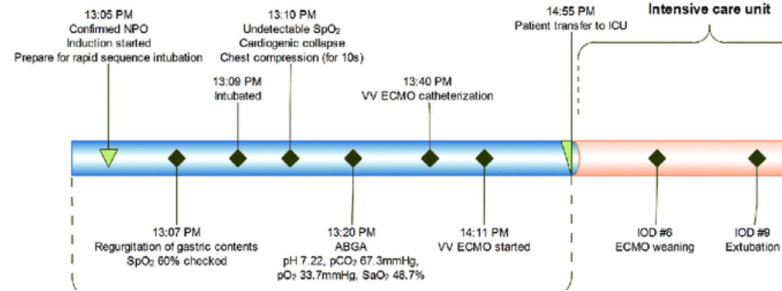


1.3.3 A donkey, a dog, a cat, and a rooster.

Multi-label classification:
Tagging




Recommendation



Sequence learning: Patient clinical event prediction




Sample Generation:
Generative Adversarial Networks



Sequence learning: Image Captioning



Sequence to sequence learning:
Translation



Unsupervised learning: self supervised learning