

# Cloud Native and Microservices

Dominic Duggan  
Stevens Institute of Technology

1

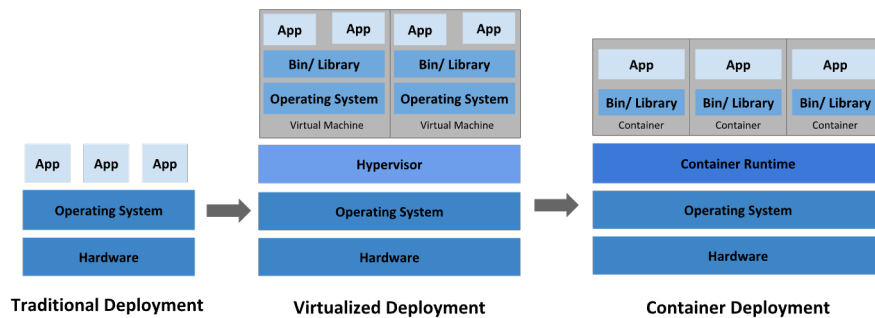
1

## **CONTAINERS ORCHESTRATION FRAMEWORKS**

2

2

## Background



3

3

## Containers

- Agile (compared to VMs)
- Continuous development, integration, deployment
- Separation of concerns
  - Create container image at build/release time
- Observability (Health)
- Environmental consistency
- Loose coupling
- Resource isolation

4

4

## Container Orchestration Frameworks

- Connect multiple containers together
- Scale up number of container instances

5

5

## Container Orchestration Frameworks

- Service discovery & load balancing
- Storage orchestration
- Automated rollouts and rollbacks
- Automated bin packing
- Self healing
- Secret and configuration management
- **Not a PaaS!**

6

6

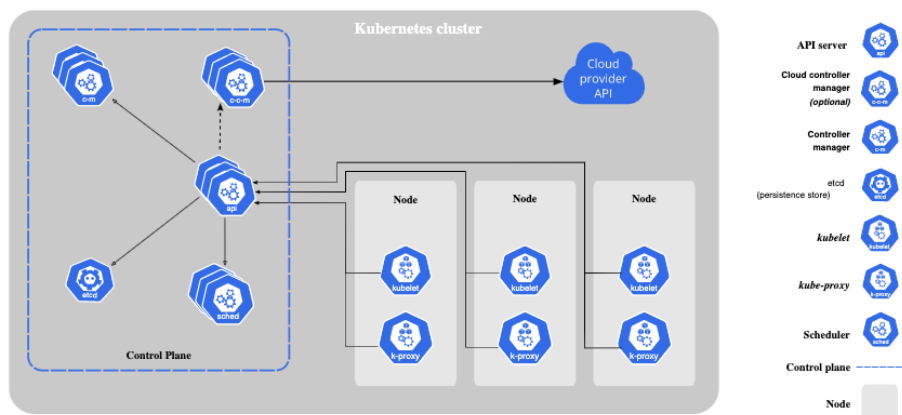
## Example: Kubernetes

- Pod: atomic workload unit (Linux container)
- Deployment: manages running pods
  - Scaling up/down
  - Rolling updates
- Service: Logical abstraction for API
  - Direct client requests to pool of pods

7

7

## Kubernetes Cluster



8

8

## Kubernetes Control Plane


- API server
  - Scales horizontally
- Etcd
  - Highly available key-value store
  - Backing store for clusters
- Schedule *pods* to run on *nodes*
- Controller managers

9

9

## Kubernetes: Node

- Worker machine
- Master schedules Pods on Nodes
- Node runs:
  - Docker runtime
  - Kubelet
    - Communication with Master
    - Manages Pods and Containers



10

10

## Node Components

- Kubelet
  - Agent that runs on each node
  - Manage containers
- Kube-proxy
  - Network proxy
  - Packet filtering
- Container runtime

11

11

## Node Status

- Addresses
  - External and internal IP address
- Conditions

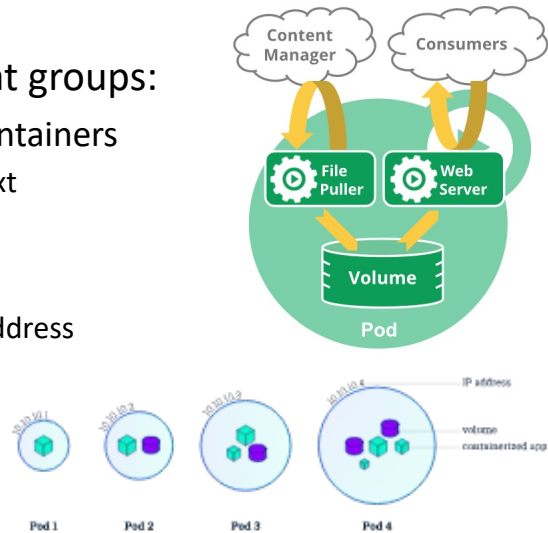
```
"conditions": [  
  {  
    "type": "Ready",  
    "status": "True",  
    "reason": "KubeletReady",  
    "message": "kubelet is posting ready status",  
    "lastHeartbeatTime": "2019-06-05T18:38:35Z",  
    "lastTransitionTime": "2019-06-05T11:41:27Z"  
  }  
]
```

12

12

# Kubernetes: Pod

- Abstraction that groups:
  - Application containers
    - Shared context
  - Resources
    - Volumes
    - Network IP address
    - Configuration



13

13

# YAML

- Yet Another Markup Language
- Basically alternative syntax for JSON
- Configuration language
  - Kubernetes, OpenStack, etc

14

14

## YAML vs JSON

### YAML

```
---
apiVersion: v1
kind: Pod
```

### JSON

```
{
  "apiVersion": "v1",
  "kind": "Pod"
}
```

15

15

## YAML vs JSON

### YAML

```
---
apiVersion: v1
kind: Pod
metadata:
  name: rss-site
labels:
  app: web
```

### JSON

```
{
  "apiVersion": "v1",
  "kind": "Pod",
  "metadata": {
    "name": "rss-site",
    "labels": {
      "app": "web"
    }
  }
}
```

**Don't use tab in YAML files!**

16

16



## YAML vs JSON

### YAML

```
args:
  - sleep
  - "1000"
  - message
  - "Bring back Saul!"
```

### JSON

```
{
  "args": [
    "sleep",
    "1000",
    "message",
    "Bring back Saul!"
  ]
}
```

17

17

## Pod in YAML

```
---
apiVersion: v1
kind: Pod
metadata:
  name: rss-site
  labels:
    app: web
spec:
  containers:
    - name: front-end
      image: nginx
      ports:
        - containerPort: 80
    - name: rss-reader
      image: rdr/php-nginx
      ports:
        - containerPort: 88
```

18

18

# Pod Template

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello
spec:
  template:
    # This is the pod template
    spec:
      containers:
      - name: hello
        image: busybox
        command: ['sh', '-c', 'echo "Hello, Kubernetes!"
                  && sleep 3600']
        restartPolicy: OnFailure
    # The pod template ends here
```

19

19

## Controllers and Workload Resources

- Controller for Resource
  - Replication and rollout
  - Automatic healing
- Workload Resources examples
  - Job
  - Deployment
  - Statefulset
  - DaemonSet

20

20

## Example Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello
spec:
  template:
    # This is the pod template
    spec:
      containers:
        - name: hello
          image: busybox:1.28
          command: ['sh', '-c', 'echo "Hello!" && sleep 3600']
          restartPolicy: OnFailure
    # The pod template ends here
```

21

21

## Kubernetes: Deployment

- Declarative update (desired state)
- Use cases
  - Ensure availability of a workload
  - Scale up for higher load
  - Change state of Pods
  - Roll back to earlier Deployment
  - Pause the Deployment
  - Check status of Deployment
  - Expose workload outside the cluster

22

22

## Example Deployment

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rss-site
spec:
  replicas: 2
  selector:
    matchLabels:
      app: web

  template:
    metadata:
      labels:
        app: web
    spec:
      containers:
        - name: front-end
          image: nginx
          ports:
            - containerPort: 80
        - name: rss-reader
          image: rdr/php-nginx
          ports:
            - containerPort: 88
```

23

23

## Example Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

24

24

# Kubernetes: CLI

```
$ kubectl get pods
$ kubectl describe pods
```

```
Name:      kubernetescamp-765bf4c7b4-2dp6k
Namespace: default
Priority:   0
Node:      minikube/172.17.0.9
Start Time: Tue, 30 Mar 2021 13:37:02 +0000
Labels:    pod-template-hash=765bf4c7b4
           run=kubernetescamp
Annotations: <none>
Status:    Running
IP:        172.18.0.6
IPs:
  IP:      172.18.0.6
Controlled By: ReplicaSet/kubernetescamp-765bf4c7b4
Containers:
  kubernetescamp:
    Container ID:  docker://859a7c163dfc451add454c06c2bc4c60928efcb60bfc772e4edfc99de4f7560e
    Image:         gcr.io/google-samples/kubernetescamp:v1
    Image ID:      docker-pullable://jocatalin/kubernetescamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f00e279c8fcc64af
    Port:          8080/TCP
    Host Port:     0/TCP
    State:         Running
    Started:       Tue, 30 Mar 2021 13:37:04 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-8k4fg (ro)
```

```
$ kubectl logs pod-name
```

25

25

# Kubernetes: CLI

```
$ kubectl exec pod-name env
```

```
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=kubernetescamp-765bf4c7b4-2dp6k
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_SERVICE_PORT=443
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
NPM_CONFIG_LOGLEVEL=info
NODE_VERSION=6.3.1
HOME=/root
```

```
$ kubectl exec --ti pod-name /bin/bash
# curl localhost:8080
```

26

26

## SERVICES

27

27

## Kubernetes Networking

- Goal: Frictionless migration from VMs to containers
- Within a pod:
  - Containers communicate via loopback (localhost)
- Pods communicate in cluster without NAT
- Service resource
  - Expose app outside cluster
  - Ingress: specific for HTTP

28

28

# Kubernetes Services

- Pods
  - Interchangeable among Pods in a ReplicaSet
  - Unique IP address
- Service
  - Discovery and routing among Pods
  - Loose coupling
  - Type:
    - ClusterIP: only within cluster
    - NodePort: each selected node in cluster
    - LoadBalancer: fixed external IP
    - ExternName: DNS

29

29

# Kubernetes Services

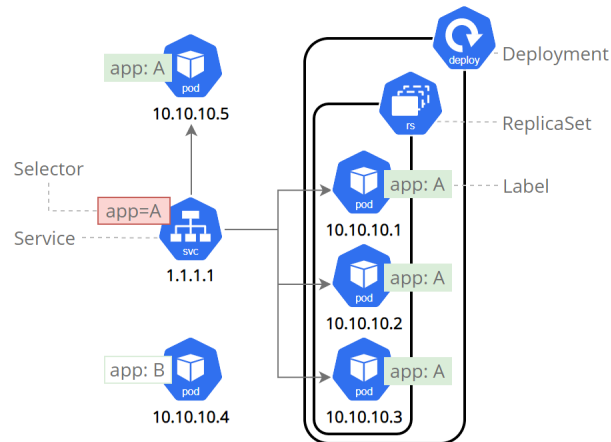
```
kind: Service
apiVersion: v1
metadata:
  name: hello-service
spec:
  selector:
    app: hello-pod
  ports:
    - port: 8080
```

30

30

# Kubernetes Services

- Match Pods using Labels and Selectors



31

31

# Kubernetes: CLI

```
$ kubectl expose deployment/service-name
--type="NodePort" --port 8080
```

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
kubernetes-bootcamp	NodePort	10.110.126.250	<none>	8080:30601/TCP

```
$ kubectl describe services/service-name
```

```
Name:
Namespace:
Labels:
Annotations:
Selector:
Type:
IP:
Port:
TargetPort:
NodePort:
Endpoints:
Session Affinity:
External Traffic Policy:
Events:
```

```
$ curl node-ip-address:30601
```

32

32



## Kubernetes: Labels

```
$ kubectl get pods -l label
$ kubectl get services -l label
$ kubectl label pod pod-name app=v1
$ kubectl describe pods pod-name
```

```
Name:      kubernetes-bootcamp-765bf4c7b4-b5vrs
Namespace: default
Priority:   0
Node:      minikube/172.17.0.72
Start Time: Tue, 30 Mar 2021 14:05:05 +0000
Labels:    app=v1
           pod-template-hash=765bf4c7b4
           run=kubernetes-bootcamp
Annotations: <none>
Status:    Running
IP:        172.18.0.2
IPs:       172.18.0.2
Controlled By: ReplicaSet/kubernetes-bootcamp-765bf4c7b4
```

```
$ kubectl get pods -l app=v1
```

33

33


## Example Deployment

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: hello-deploy
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: hello-pod
    spec:
      containers:
        - name: hello-container
          image: ex.com/hello-cloud:1
          imagePullPolicy: IfNotPresent
          livenessProbe:
            httpGet:
              path: /
              port: 8080
          readinessProbe:
            httpGet:
              path: /resources/hello
              port: 8080
          restartPolicy: Always
```

34

34

# Kubernetes: Scaling




```
kubectl scale deployments/name --replicas=4
```

35

35

# Kubernetes: Rolling Updates



```
kubectl scale deployments/name --replicas=4
```

36

36

## Exposing External IP Address

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/name: load-balancer-example
  name: hello-world
spec:
  replicas: 5
  selector:
    matchLabels:
      app.kubernetes.io/name: load-balancer-example
  template:
    metadata:
      labels:
        app.kubernetes.io/name: load-balancer-example
    spec:
      containers:
        - image: gcr.io/google-samples/node-hello:1.0
          name: hello-world
          ports:
            - containerPort: 8080
```

37

37

## Exposing External IP Address

- Deploy  
kubect1 apply -f load-balance-example.yaml
- Expose with a Service  
kubect1 expose deployment hello-world  
--type=LoadBalancer --name=my-service  
kubect1 get services my-service
- Output

NAME	TYPE	EXTERNAL-IP	PORT(S)	AGE
my-service	LoadBalancer	104.198.205.71	8080/TCP	54s

38

38

## Exposing External IP Address

- Detailed description

```
kubectl describe services my-service
```


- Output

```
Name:          my-service
Namespace:     default
Labels:        app.kubernetes.io/name=load-balancer-example
Annotations:   <none>
Selector:      app.kubernetes.io/name=load-balancer-example
Type:          LoadBalancer
IP:            10.3.245.137
LoadBalancer Ingress: 104.198.205.71
Port:          <unset> 8080/TCP
NodePort:      <unset> 32377/TCP
Endpoints:     10.0.0.6:8080,10.0.1.6:8080,10.0.1.7:8080 +
2 more...
Session Affinity: None
Events:        <none>
```

39

39

## Kubernetes Ingress



40

40

# Kubernetes Ingress

```
kind: Ingress
apiVersion: extensions/v1beta1
metadata:
  name: hello-ingress
spec:
  rules:
  - host: hello.example.com
    http:
      paths:
      - path: /
        backend:
          serviceName: hello-service
          servicePort: 8080
```

41

41

# ConfigMap

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: hello-cloud-config
data:
  application.properties: |
    hello.greeting=Hello from Kubernetes
    hello.name=Java EE
```

42

42

## Deployment

```
kind: Deployment
spec:
  replicas: 1
  template:
    spec:
      containers:
        - name: hello-container
          volumeMounts:
            - name: config-volume
              mountPath: /opt/config
      volumes:
        - name: config-volume
          configMap:
            name: hello-cloud-config
```

43

43

## Connecting External Services

```
@ApplicationScoped
public class HelloCloudProcessor {
    private Client client;
    private WebTarget target;

    @PostConstruct
    private void initClient() {
        client = ClientBuilder...
        target = client.target("http://cloudprocessor:
                               8080/processor/resources/hello");
    }

    public String processGreeting() {
        ...
    }
}
```

44

44

## Configuring Orchestrated Applications

```
public class HelloGreeter {  
  
    @Inject @Config("hello.greeting")  
    String greeting;  
  
    @Inject @Config("hello.name")  
    String greetingName;  
  
    public String processGreeting() {  
        return greeting + ", " + greetingName;  
    }  
}
```

45

45

## Configuring Orchestrated Applications

```
@ApplicationScoped  
public class ConfigurationExposer {  
  
    private final Properties properties = new Properties();  
  
    @PostConstruct  
    private void initProperties() {  
        try (InputStream inputStream =  
            new FileInputStream  
                ("/opt/config/application.properties")) {  
            properties.load(inputStream);  
        } catch (IOException e) {  
            throw new IllegalStateException("...", e);  
        }  
    }  
}
```

46

46

## Configuring Orchestrated Applications

```
@ApplicationScoped
public class ConfigurationExposer {

    private final Properties properties = new Properties();

    @Produces @Config("")
    public String exposeConfig
        (InjectionPoint injectionPoint) {
        Config config =
            injectionPoint.getAnnotated()
                .getAnnotation(Config.class);
        if (config != null)
            return properties.getProperty(config.value());
        return null;
    }
}
```

47

47