

NoSQL Data stores

Dominic Duggan

Stevens Institute of Technology

Based on materials by A. Haeberlen, Z. Ives, E.Meijer

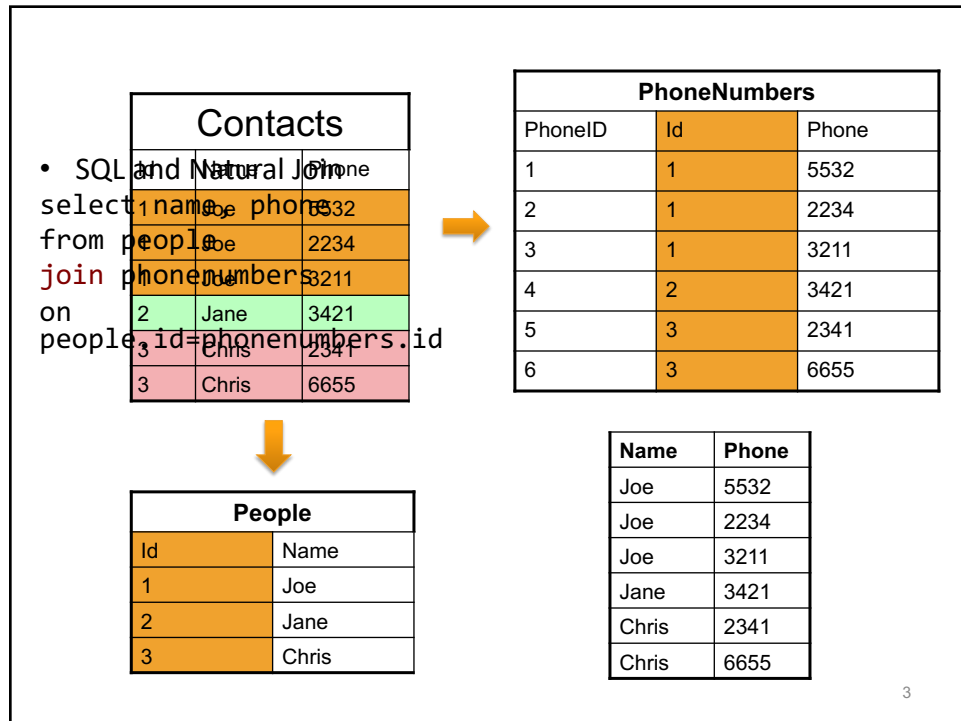
1

1

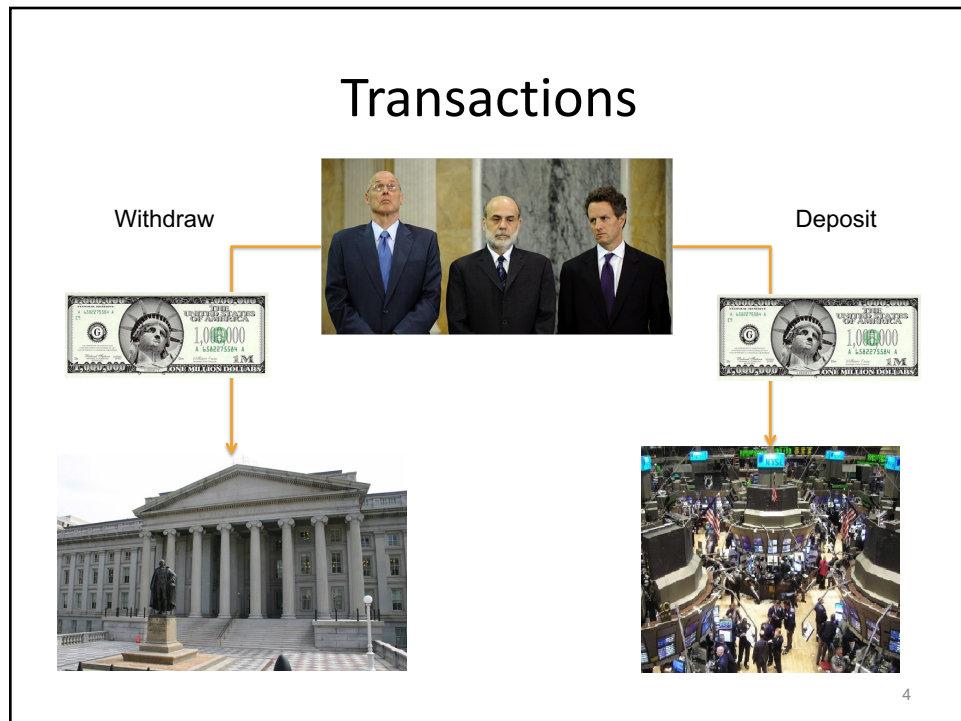
DATABASE REVIEW

2

2



3



4

Transactions

Withdraw



Deposit



5

5

Transactions

Withdraw



Deposit



6

6

Transactional Operations

amazon.com[®]

```
Begin Transaction
  AmtOwed = Amazon.Read (CustID);
  Balance = Paypal.Read (AccountID);
  if (AmtOwed ≤ Balance) {
    Paypal.Withdraw (AccountID, AmtOwed);
    Amazon.Deposit (CustID, AmtOwed);
  } else {
    Abort Transaction;
  }
End Transaction
```

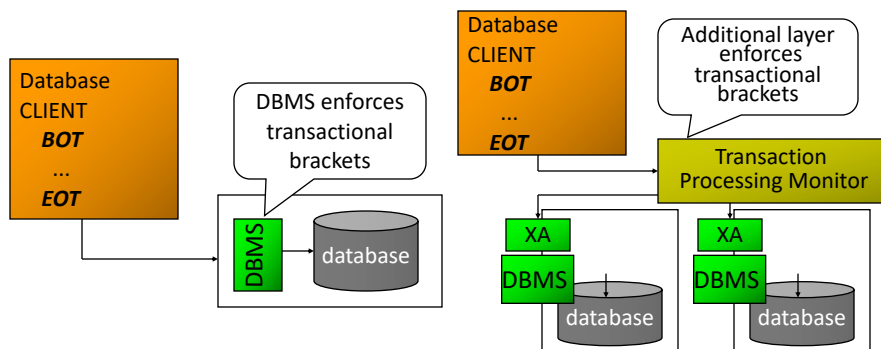
PayPal[®]

7

7

Distributed Databases

- Transactions and atomic commitment



8

8

Relational Database Summary


- Database Schema
 - Normalized for efficiency
- SQL for ad-hoc queries
- Transactional updates
 - Attomic
 - Consistent
 - Isolated
 - Durable

9

9

Challenge: Big Data™

- Historical approach:
vertical scaling
 - Limited
- Modern approach:
horizontal scaling
 - *Sharding*
 - Azure: Federated SQL databases
 - Applications see data partitioning
 - No joins across partitions



10

10

SQL vs NoSQL

Relational

- Database Schema
 - Business data model
- SQL for ad-hoc queries
- ACID properties
 - Atomic
 - Consistent
 - Isolated
 - Durable

NoSQL

- Unstructured
 - Web server logs
- Map-Reduce
- BASE properties
 - Basically Available
 - Soft state
 - Eventually consistent

11

11

Fault-tolerance



12

12

NOSQL DATA MODELS

13

13

NoSQL taxonomy

- Key-Value stores (DHT)
- Column stores
- Document stores


14

14

- Key-Value stores (DHT) Amazon Dynamo
- Column stores Google Bigtable, Cassandra
- Document stores CouchDB, MongoDB, SimpleDB

15

Key / Value




```
{
  "name": "uri",
  "ssn": "213445",
  "hobbies": [ "H", "H", "H" ],
  ...
}
```

```
{
  { ... }
}
```

```
{
  { ... }
}
```

16

Key/Value

- Have the key? Get the value
 - Map/Reduce (sometimes)
 - Good for
 - cache aside (e.g. Hibernate 2nd level cache)
 - Simple, id based interactions (e.g. user profiles)
- In most cases, values are opaque


K1	V1
K2	V2
K3	V3
K4	V1

17


17

Key/Value

- Scaling out is relatively easy
 - just hash the keys
- Fixed vs. consistent hashing



18




18

Column Based

- Mostly derived from Google's BigTable papers
- One giant table of rows and columns
 - Column == pair (name and a value, sometimes timestamp)
 - Table is sparse:

$$(\text{\#rows}) \times (\text{\#columns}) \geq (\text{\#values})$$



19

19

Column Based

- Query on row key
 - Or column value (aka secondary index)
- Good for a constantly changing, (albeit flat) domain model

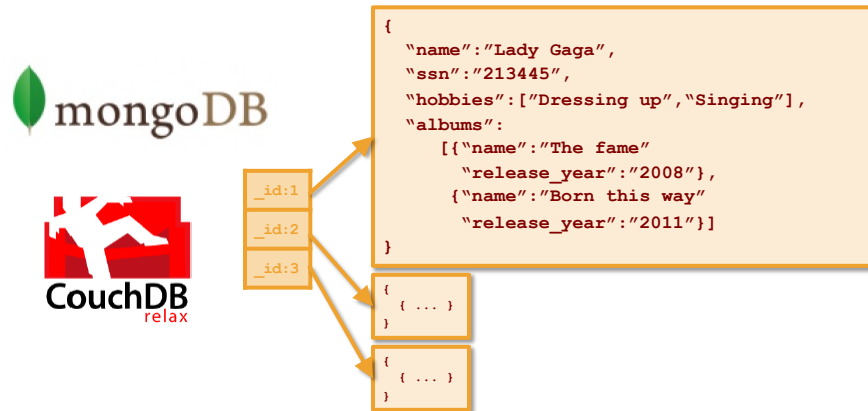


20

20

Document

- Think JSON (or BSON, or XML)



21

21

Document

- Model is not flat, data store is aware of it
 - Arrays, nested documents
- Better support for ad hoc queries
 - MongoDB excels at this
- Very intuitive model
- Flexible schema

```
> db.people.find({age: {$gt: 27}})
{ "_id" : ObjectId("4bed80b20b4acd070c593bac"), "name" : "John", "age" : 28 }
{ "_id" : ObjectId("4bed80b0b4acd070c593bad"), "name" : "Steve", "age" : 29 }
```

22

22