

NOSQL DATA MODELS

2

2

NoSQL taxonomy

- Key-Value stores
- Column stores
- Document stores

3

3

NoSQL taxonomy

- Key-Value stores Amazon Dynamo
- Column stores Google Bigtable,
Cassandra
- Document stores CouchDB, MongoDB,
SimpleDB

4

- 4

A Universe of Data Models

The diagram illustrates three data models: Key/Value, Column, and Document.

- Key / Value:** Represented by a 4x2 grid of orange squares, indicating a simple key-value pair structure.
- Column:** Represented by a 4x4 grid of orange squares, with some squares missing (white), indicating a columnar structure.
- Document:** Represented by a 3x1 grid of orange squares, with arrows pointing to a large green box containing a JSON document and two smaller green boxes containing fragments of the document, indicating a document-oriented structure.

```
{
  "name": "uri",
  "ssn": "213445",
  "hobbies": ["...", "..."],
  "...": {
    "...": "...",
    "...": "...",
    "...": "..."
  }
}
```

```
{ ... }
```

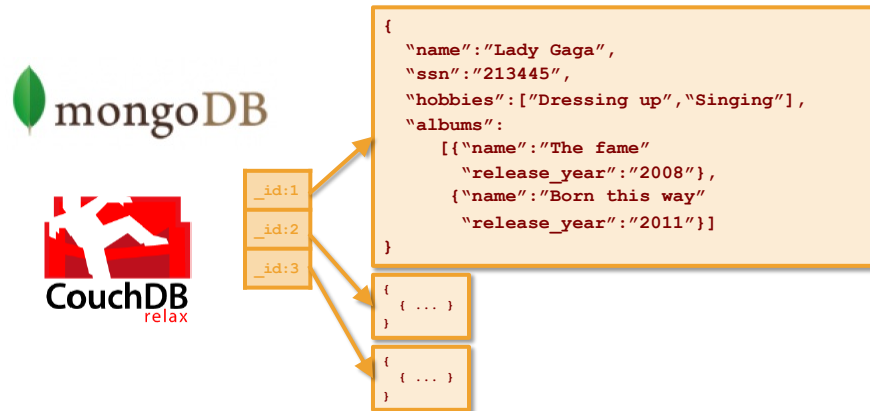
```
{ ... }
```

5



Document

- Think JSON (or BSON, or XML)



6

6

Key/Value

- Have the key? Get the value
 - Map/Reduce (sometimes)
 - Good for
 - cache aside (e.g. Hibernate 2nd level cache)
 - Simple, id based interactions (e.g. user profiles)
- In most cases, values are opaque


K1	V1
K2	V2
K3	V3
K4	V1

7

7

Column Based

- Mostly derived from Google's BigTable papers
- One giant table of rows and columns
 - Column == pair (name and a value, sometimes timestamp)
 - Table is sparse:
 $(\text{\#rows}) \times (\text{\#columns}) \geq (\text{\#values})$



8

8

Column Based

- Query on row key
 - Or column value (aka secondary index)
- Good for a constantly changing, (albeit flat) domain model



9

9



BIGTABLE: COLUMN-BASED STORE

10

10

Data Model


A single column

Name	colA	value1	Value
------	------	--------	-------

11

11

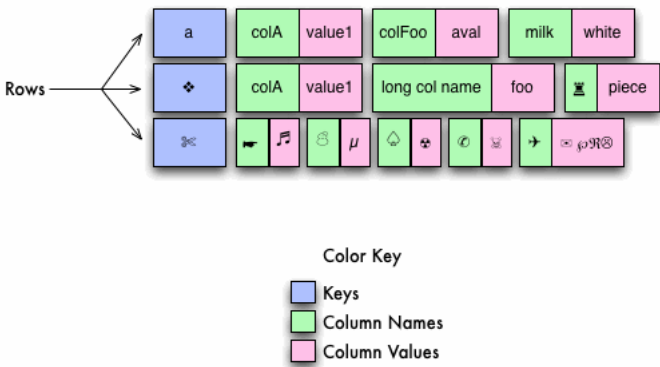
Data Model



12

12

Data Model



13

13

Data Model - Vocabulary

- **Keyspace** – like namespace for unique keys.
- **Column Family** – very much like a table... but not quite.
- **Key** – a key that represents row (of columns)
- **Column** – representation of value with:
 - Column name
 - Value
 - Timestamp

14

14

Data Model - Column Family

- Similar to SQL tables
- Groups of column keys
- Hundreds of static column families
- Syntax is family:key, e.g., Language:English, Language:German, etc

15

15

Data Model - Rows

- Primary key for objects
- All keys are arbitrary length binaries

Users:	CF
foo:	ROW
emailAddress: foo@bar.com,	COLUMN
webSite: http://bar.com	COLUMN
bar:	ROW
emailAddress: bug@example.org	COLUMN
Stats:	CF
foo:	ROW
visits: 243	COLUMN

16

16

Data Model Summary


- Keyspace
 - ColumnFamily
 - Row (indexed)
 - Key
 - Columns
 - Name (sorted)
 - Value



17

Bigtable

- A sparse, distributed persistent multi-dimensional sorted map.
- The map is indexed by a row key, column key, and a timestamp.
- Output value is array of bytes.
 - (row: byte[], column: byte[], time: int64) → byte[]



18

Rows

- Read/write of data under a single row is atomic
- Data ordered by row key
- Row range dynamically partitioned
- **Tablet:** partition (row range)
 - Unit of distribution and load-balancing
- Objective: make read operations single-sited!
 - E.g., In Webtable: com.google.maps instead of maps.google.com.

19

Timestamps

- 64 bit integers
- Items in a cell stored in decreasing timestamp order
- Application specifies how many versions (n) of data items are maintained in a cell.
 - Bigtable garbage collects obsolete versions.

20

20

Application 1: Google Analytics

- Webmaster stats:
 - Number of unique visitors per day
 - Page views per URL per day,
 - Percentage of users that made a purchase given that they earlier viewed a specific page.
- How?
 - JavaScript embedded in web pages, records:
 - User identifier
 - The page being fetched

21

21

Application 1: Google Analytics

- Raw click table (~ 200 TB)
 - A row for each end-user session
 - Row name includes website name and time of session
 - Clustering of sessions that visit the same web site
- Summary table (~ 20 TB)
 - Predefined summaries for each web site.
 - Generated from raw click table by periodic MapReduce jobs.
 - Row name includes website's name
 - Column family is aggregate summaries

22

22

Application 2: Personalized Search

- Records user queries and clicks across Google
- Users browse their search histories
- Personalized search results
- One Bigtable:
 - Row name is userid
 - Column family for each action type, e.g., web queries, clicks.
 - User profiles are generated using MapReduce.
 - Replicated geographically to reduce latency and increase availability.

23

23