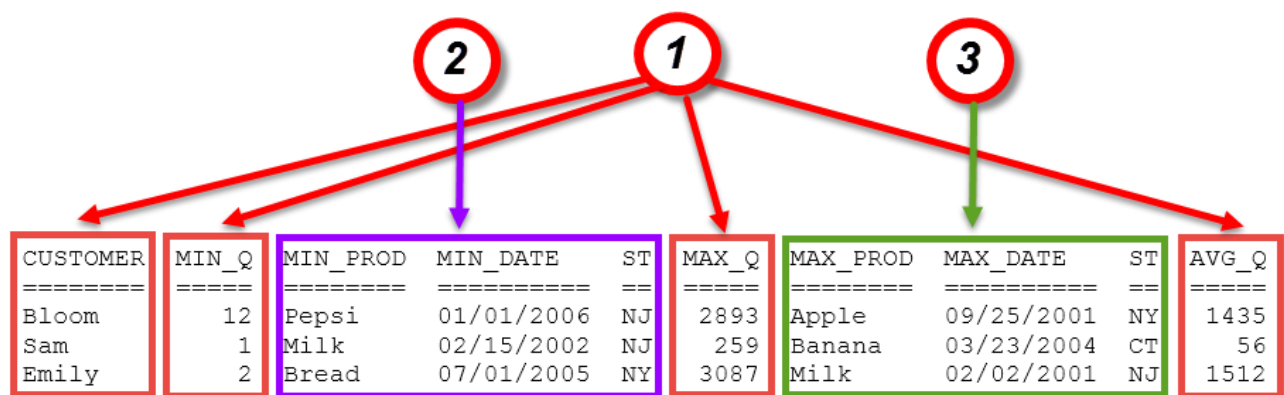


CS561 – Programming Assignment 1

“The Idea”

To help you understand the mechanics of writing the SQL queries for the assignment, I have outlined below the “structure” of the query results, and how each part is produced and how the three parts are combined to produce the final result set.

We will discuss this further in the next lecture, so please read it carefully and be prepared to participate in discussions.



As shown in the diagram above, the output of the query #1 consists of 3 main parts, and they are constructed in the order listed (query #2 works in a very similar fashion):

1. This is the easy part, where you are simply computing MIN_Q (minimum quantity), MAX_Q (maximum quantity) and AVG_Q (average quantity) based on the grouping attribute, CUSTOMER. It is a simple “group-by” query (or “aggregation operator” in relational algebra).
2. Next step is to capture the corresponding information for each pair of (CUSTOMER, MIN_Q) – that is, MIN_PROD, MIN_DATE and ST (that is, the corresponding PRODUCT, DATE and STATE for the given (CUSTOMER, MIN_Q)). The “challenge” is that the corresponding information is only available in the SALES table; therefore, you will need to “join” the result of Part #1 with the SALES table to pull the corresponding information from the SALES table.

As you can see below, if you join the results of Part #1 with the SALES table using the condition (predicate) of “**part1.CUSTOMER = sales.CUST and part1.MIN_Q = sales.QUANT**”, you will find the corresponding information (highlighted in **YELLOW**, with a PURPLE box around them) in the SALES table (NOTE: the data you see in the diagrams below are made up and do not match the actual contents of the SALES table you have – this is just for illustrative purposes).

CUSTOMER	MIN_Q	MIN_PROD	MIN_DATE	ST		CUST	PROD	DAY	MONTH	YEAR	STATE	QUANT
Bloom	12	Pepsi	01/01/2006	NJ		Bloom	Pepsi	2	12	2001	NY	4232
Sam	1	Milk	02/15/2002	NJ		Knuth	Bread	23	5	2005	PA	4167
Emily	2	Bread	07/01/2005	NY		Emily	Pepsi	22	1	2006	CT	4404
						Emily	Fruits	11	1	2000	NJ	4369
						Helen	Milk	7	11	2006	CT	210
						Emily	Soap	2	4	2002	CT	2549
						Bloom	Pepsi	1	1	2006	NJ	12
						Bloom	Yogurt	25	7	2004	PA	17
						Helen	Pepsi	14	3	2002	NJ	3891
						Emily	Bread	28	9	2005	PA	42
						Sam	Cookies	20	11	2004	NY	3376
						Knuth	Milk	5	2	2007	PA	126
						Helen	Coke	11	4	2001	NY	668
						Emily	Butter	5	7	2005	NJ	3840
						Emily	Yogurt	7	10	2005	NY	730
						Sam	Soap	12	2	2001	NJ	165
						Knuth	Coke	6	1	2003	CT	1557
						Sam	Milk	9	8	2001	NY	1132
												4001

Results of Part #1

SALES table

- Next step is to capture the corresponding information for each pair of (CUSTOMER, MAX_Q) – that is, MAX_PROD, MAX_DATE and ST (that is, the corresponding PRODUCT, DATE and STATE for the given (CUSTOMER, MAX_Q), and the idea is exactly the same as the Part #2.

Things to keep in mind:

- The result of each part needs to be “saved” (or stored) in a temporary result set (e.g., a table, view, etc.) – for this, the best method to use the “WITH” clause. WITH is a syntactic construct that allows you to define a “temporary view” (a derived relation/table), and you can find the details in slides 3.43 and 3.44. You can also use “VIEWS” (3.45 – 3.47); however, it’s a bit cumbersome to use VIEWS (as compared to WITH clauses). Also note that the WITH syntax for PostgreSQL is slightly different from the one presented in the slides (e.g., in PostgreSQL, you don’t need to list the column names when you’re defining the temporary results).
- The three parts of the query results are built on top of the results of the previous Parts – that is, the results of Part #2 is created using the results of Part #1, and similarly, the results of Part #3 is created based on the results of Part #3.
- For the date columns of the result set (output), you don’t need to concatenate the month, day and year into a date column (e.g., 01/01/2016; instead, just project MONTH, DAY and YEAR columns separately, e.g., 1 | 1 | 2016 or directly use the “date” column.